

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Problem Statement

A retail company "ABC Private Limited" wants to understand the customer purchase behaviour (specifically, purchase amount) against various products of different categories. They have shared purchase summary of various customers for selected high volume products from last month. The data set also contains customer demographics (age, gender, marital status, city_type, stay_in_current_city), product details (product_id and product category) and Total purchase_amount from last month.

Now, they want to build a model to predict the purchase amount of customer against various products which will help them to create personalized offer for customers against different products.

```
In [2]: #importing the train data
df_train = pd.read_csv('train.csv')
```

```
In [3]: df_train.head(3)
```

```
Out[3]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1000001	P00069042	F	0-17	10	A	2	0
1	1000001	P00248942	F	0-17	10	A	2	0
2	1000001	P00087842	F	0-17	10	A	2	0

```
In [4]: df_train.shape
```

```
Out[4]: (550068, 12)
```

```
In [5]: #importing the test data
df_test = pd.read_csv('test.csv')
```

```
In [6]: df_test.head(3)
```

Out[6]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1000004	P00128942	M	46-50	7	B	2	
1	1000009	P00113442	M	26-35	17	C	0	
2	1000010	P00288442	F	36-45	1	B	4+	

In [7]: `df_test.shape`

Out[7]: (233599, 11)

In [8]: `#append both train and test data`
`df = df_train.append(df_test)`

C:\Users\vcyad\AppData\Local\Temp\ipykernel_11136\4095870650.py:2: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
`df = df_train.append(df_test)`

In [9]: `df.head()`

Out[9]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	1000001	P00069042	F	0-17	10	A	2	
1	1000001	P00248942	F	0-17	10	A	2	
2	1000001	P00087842	F	0-17	10	A	2	
3	1000001	P00085442	F	0-17	10	A	2	
4	1000002	P00285442	M	55+	16	C	4+	

In [10]: `df.shape`

Out[10]: (783667, 12)

Understanding the data

In [11]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               783667 non-null  int64
1   Product_ID                            783667 non-null  object
2   Gender                                783667 non-null  object
3   Age                                    783667 non-null  object
4   Occupation                             783667 non-null  int64
5   City_Category                          783667 non-null  object
6   Stay_In_Current_City_Years            783667 non-null  object
7   Marital_Status                         783667 non-null  int64
8   Product_Category_1                     783667 non-null  int64
9   Product_Category_2                     537685 non-null  float64
10  Product_Category_3                     237858 non-null  float64
11  Purchase                               550068 non-null  float64
dtypes: float64(3), int64(4), object(5)
memory usage: 77.7+ MB
```

In [12]:

df.describe()

Out[12]:

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product
count	7.836670e+05	783667.000000	783667.000000	783667.000000	537685.000000	23
mean	1.003029e+06	8.079300	0.409777	5.366196	9.844506	
std	1.727267e+03	6.522206	0.491793	3.878160	5.089093	
min	1.000001e+06	0.000000	0.000000	1.000000	2.000000	
25%	1.001519e+06	2.000000	0.000000	1.000000	5.000000	
50%	1.003075e+06	7.000000	0.000000	5.000000	9.000000	
75%	1.004478e+06	14.000000	1.000000	8.000000	15.000000	
max	1.006040e+06	20.000000	1.000000	20.000000	18.000000	

In [13]:

df.drop(['User_ID'], axis = 1, inplace = True)

In [14]:

df.head(3)

Out[14]:

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Pr
0	P00069042	F	0-17	10	A	2	0	
1	P00248942	F	0-17	10	A	2	0	
2	P00087842	F	0-17	10	A	2	0	

In [15]:

#df['Gender'] = pd.get_dummies(df['Gender'], drop_first = 1)

```
In [16]: #handling categorical feature gender using map function
df['Gender'] = df['Gender'].map({'F':0, 'M':1})
df.head()
```

```
Out[16]:
```

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Pr
0	P00069042	0	0-17	10	A	2	0	
1	P00248942	0	0-17	10	A	2	0	
2	P00087842	0	0-17	10	A	2	0	
3	P00085442	0	0-17	10	A	2	0	
4	P00285442	1	55+	16	C	4+	0	

```
In [17]: #handling categorical feature age
df['Age'].unique()
```

```
Out[17]: array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
      dtype=object)
```

```
In [18]: df['Age']=df['Age'].map({'0-17':1, '18-25':2, '26-35':3, '36-45':4, '46-50':5, '51-55':6, '55+':7})
```

second technique

from sklearn import preprocessing

label_encoder object knows how to understand word labels.

```
label_encoder = preprocessing.LabelEncoder()
```

Encode labels in column 'species'

```
df['Age'] = label_encoder.fit_transform(df['Age'])
```

```
df['Age'].unique()
```

```
In [19]: df.head()
```

Out[19]:

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Pr
0	P00069042	0	1	10	A	2	0	
1	P00248942	0	1	10	A	2	0	
2	P00087842	0	1	10	A	2	0	
3	P00085442	0	1	10	A	2	0	
4	P00285442	1	7	16	C	4+	0	

In [20]:

```
print(df['City_Category'].unique(),'\n')
print(df['City_Category'].value_counts())
```

```
['A' 'C' 'B']

B    329739
C    243684
A    210244
Name: City_Category, dtype: int64
```

In [21]:

```
pd.get_dummies(df['City_Category'])
```

Out[21]:

	A	B	C
0	1	0	0
1	1	0	0
2	1	0	0
3	1	0	0
4	0	0	1
...
233594	0	1	0
233595	0	1	0
233596	0	1	0
233597	0	0	1
233598	0	1	0

783667 rows × 3 columns

In [22]:

```
df_city = pd.get_dummies(df['City_Category'],drop_first = True)
```

In [23]:

```
df_city.head()
```

Out[23]:

	B	C
0	0	0
1	0	0
2	0	0
3	0	0
4	0	1

```
In [24]: df = pd.concat([df,df_city], axis = 1)
df
```

Out[24]:

	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
0	P00069042	0	1	10	A	2	
1	P00248942	0	1	10	A	2	
2	P00087842	0	1	10	A	2	
3	P00085442	0	1	10	A	2	
4	P00285442	1	7	16	C	4+	
...
233594	P00118942	0	3	15	B	4+	
233595	P00254642	0	3	15	B	4+	
233596	P00031842	0	3	15	B	4+	
233597	P00124742	0	5	1	C	4+	
233598	P00316642	0	5	0	B	4+	

783667 rows × 13 columns

```
In [25]: #drop city category feature
df.drop('City_Category', axis = 1, inplace = True)
df.head()
```

Out[25]:

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category
0	P00069042	0	1	10	2	0	
1	P00248942	0	1	10	2	0	
2	P00087842	0	1	10	2	0	
3	P00085442	0	1	10	2	0	
4	P00285442	1	7	16	4+	0	

```
In [26]: #missing values
df.isnull().sum()
```

```
Out[26]: Product_ID      0
        Gender        0
        Age           0
        Occupation    0
        Stay_In_Current_City_Years  0
        Marital_Status  0
        Product_Category_1  0
        Product_Category_2  245982
        Product_Category_3  545809
        Purchase      233599
        B             0
        C             0
        dtype: int64
```

```
In [27]: # focusing on replacing missing values
        df['Product_Category_2'].unique()
```

```
Out[27]: array([nan,  6., 14.,  2.,  8., 15., 16., 11.,  5.,  3.,  4., 12.,  9.,
        10., 17., 13.,  7., 18.])
```

```
In [28]: df['Product_Category_2'].value_counts()
```

```
Out[28]: 8.0      91317
        14.0     78834
        2.0      70498
        16.0     61687
        15.0     54114
        5.0      37165
        4.0      36705
        6.0      23575
        11.0     20230
        17.0     19104
        13.0     15054
        9.0       8177
        12.0      7801
        10.0      4420
        3.0       4123
        18.0      4027
        7.0        854
        Name: Product_Category_2, dtype: int64
```

```
In [29]: df['Product_Category_2'].mode()[0]
```

```
Out[29]: 8.0
```

```
In [30]: #replace the missing values with mode
        df['Product_Category_2']=df['Product_Category_2'].fillna(df['Product_Category_2'].mode[0])
```

```
In [31]: df['Product_Category_2'].isnull().sum()
```

```
Out[31]: 0
```

```
In [32]: #product_category_3 replace missing values
        df['Product_Category_3'].unique()
```

```
Out[32]: array([nan, 14., 17.,  5.,  4., 16., 15.,  8.,  9., 13.,  6., 12.,  3.,
        18., 11., 10.])
```

```
In [33]: df['Product_Category_3'].value_counts()
```

```
Out[33]: 16.0    46469
         15.0    39968
         14.0    26283
         17.0    23818
         5.0     23799
         8.0     17861
         9.0     16532
         12.0    13115
         13.0     7849
         6.0     6888
         18.0     6621
         4.0     2691
         11.0     2585
         10.0     2501
         3.0       878
         Name: Product_Category_3, dtype: int64
```

```
In [34]: #replace missing values with mode
df['Product_Category_3']=df['Product_Category_3'].fillna(df['Product_Category_3'].mode
```

```
In [35]: df.head()
```

```
Out[35]:
```

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category
0	P00069042	0	1	10	2	0	
1	P00248942	0	1	10	2	0	
2	P00087842	0	1	10	2	0	
3	P00085442	0	1	10	2	0	
4	P00285442	1	7	16	4+	0	

```
In [36]: df.shape
```

```
Out[36]: (783667, 12)
```

```
In [37]: df['Stay_In_Current_City_Years'].unique()
```

```
Out[37]: array(['2', '4+', '3', '1', '0'], dtype=object)
```

```
In [38]: df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].str.replace('+','')
```

C:\Users\vcyad\AppData\Local\Temp\ipykernel_11136\2063355665.py:1: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].str.replace
('+','')
```

```
In [39]: df.head()
```


Out[39]:

	Product_ID	Gender	Age	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category
0	P00069042	0	1	10	2	0	
1	P00248942	0	1	10	2	0	
2	P00087842	0	1	10	2	0	
3	P00085442	0	1	10	2	0	
4	P00285442	1	7	16	4	0	

In [40]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Product_ID                            783667 non-null object
1   Gender                                783667 non-null int64
2   Age                                    783667 non-null int64
3   Occupation                            783667 non-null int64
4   Stay_In_Current_City_Years            783667 non-null object
5   Marital_Status                        783667 non-null int64
6   Product_Category_1                    783667 non-null int64
7   Product_Category_2                    783667 non-null float64
8   Product_Category_3                    783667 non-null float64
9   Purchase                              550068 non-null float64
10  B                                      783667 non-null uint8
11  C                                      783667 non-null uint8
dtypes: float64(3), int64(5), object(2), uint8(2)
memory usage: 67.3+ MB
```

In [41]: `#convert object into integers`
`df['Stay_In_Current_City_Years'] = df['Stay_In_Current_City_Years'].astype(int)`
`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Product_ID                            783667 non-null object
1   Gender                                783667 non-null int64
2   Age                                    783667 non-null int64
3   Occupation                            783667 non-null int64
4   Stay_In_Current_City_Years            783667 non-null int32
5   Marital_Status                        783667 non-null int64
6   Product_Category_1                    783667 non-null int64
7   Product_Category_2                    783667 non-null float64
8   Product_Category_3                    783667 non-null float64
9   Purchase                              550068 non-null float64
10  B                                      783667 non-null uint8
11  C                                      783667 non-null uint8
dtypes: float64(3), int32(1), int64(5), object(1), uint8(2)
memory usage: 64.3+ MB
```

```
In [42]: df['B'] = df['B'].astype(int)
df['C'] = df['C'].astype(int)
```

```
In [43]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 783667 entries, 0 to 233598
Data columns (total 12 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   Product_ID                             783667 non-null  object
 1   Gender                                 783667 non-null  int64
 2   Age                                    783667 non-null  int64
 3   Occupation                             783667 non-null  int64
 4   Stay_In_Current_City_Years             783667 non-null  int32
 5   Marital_Status                         783667 non-null  int64
 6   Product_Category_1                     783667 non-null  int64
 7   Product_Category_2                     783667 non-null  float64
 8   Product_Category_3                     783667 non-null  float64
 9   Purchase                               550068 non-null  float64
10   B                                       783667 non-null  int32
11   C                                       783667 non-null  int32
dtypes: float64(3), int32(3), int64(5), object(1)
memory usage: 68.8+ MB
```

```
In [44]: #visualizing age vs purchased
sns.barplot('Age', 'Purchase', hue = 'Gender', data = df)
```

C:\Users\vcyad\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
<AxesSubplot:xlabel='Age', ylabel='Purchase'>
```

```
Out[44]:
```




Observation: purchasing of men is high than women

```
In [45]: #visualizing purchase with occupation
sns.barplot('Occupation', 'Purchase', hue = 'Gender', data = df)
```

C:\Users\vcyad\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[45]: <AxesSubplot:xlabel='Occupation', ylabel='Purchase'>




In [46]: `sns.barplot('Product_Category_1', 'Purchase', hue='Gender', data=df)`

C:\Users\vcyad\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[46]: <AxesSubplot:xlabel='Product_Category_1', ylabel='Purchase'>




In [47]: `sns.barplot('Product_Category_2', 'Purchase', hue='Gender', data=df)`

C:\Users\vcyad\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[47]: <AxesSubplot:xlabel='Product_Category_2', ylabel='Purchase'>



```
In [48]: sns.barplot('Product_Category_3', 'Purchase', hue='Gender', data=df)
```

C:\Users\vcyad\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[48]: <AxesSubplot:xlabel='Product_Category_3', ylabel='Purchase'>
```

